

**Problem 1**

This question should be answered using the **Default** data set. In Chapter 4 on classification, we used logistic regression to predict the probability of **default** using **income** and **balance**. Now we will estimate the test error of this logistic regression model using the validation set approach. Do not forget to set a random seed before beginning your analysis.

```
> library(ISLR)
```

**(a) Fit a logistic regression model that predicts **default** using **income** and **balance**.**

```
> set.seed(1)
```

```
> fit.glm = glm(default ~ income+balance, family = "binomial", data = Default)
```

```
> summary(fit.glm)
```

Call:

```
glm(formula = default ~ income + balance, family = "binomial",
    data = Default)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4725	-0.1444	-0.0574	-0.0211	3.7245

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.154e+01	4.348e-01	-26.545	< 2e-16 ***
income	2.081e-05	4.985e-06	4.174	2.99e-05 ***
balance	5.647e-03	2.274e-04	24.836	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2920.6 on 9999 degrees of freedom

Residual deviance: 1579.0 on 9997 degrees of freedom

AIC: 1585

Number of Fisher Scoring iterations: 8

**(b) Using the validation set approach, estimate the test error of this model. You need to perform the following steps:**

**i. Split the sample set into a training set and a validation set.**

```
> train = sample(dim(Default)[1], dim(Default)[1]/2)
```

ii. Fit a logistic regression model using only the training data set.

```
> fit.glm = glm(default ~ income+balance, family = "binomial", data = Default, subset = train)
> summary(fit.glm)
```

Call:

```
glm(formula = default ~ income + balance, family = "binomial",  
    data = Default, subset = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.5830	-0.1428	-0.0573	-0.0213	3.3395

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.194e+01	6.178e-01	-19.333	< 2e-16 ***
income	3.262e-05	7.024e-06	4.644	3.41e-06 ***
balance	5.689e-03	3.158e-04	18.014	< 2e-16 ***

— — —

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1523.8 on 4999 degrees of freedom

Residual deviance: 803.3 on 4997 degrees of freedom

AIC: 809.3

Number of Fisher Scoring iterations: 8

iii. Obtain a prediction of default status for each individual in the validation set using a threshold of 0.5.

```
> glm_prob = predict(fit.glm, newdata = Default[-train, ], type = "response")
```

```
> glm_pred = rep("No",5000)
```

```
> glm pred[glm prob > 0.5] = "Yes"
```

```
> glm pred
```

[1] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"

[37] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"

[73] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "Yes" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"

[109] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"

[145] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "Yes" "No"

[illegible][illegible]

[253] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"  
"No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"

[illegible][illegible][illegible][illegible]

```
[433] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[469] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[505] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[541] "No" "Yes" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[577] "Yes" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[613] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "Yes" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[649] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[685] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[721] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "Yes" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[757] "No" "No" "No" "No" "No" "No" "Yes" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[793] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "Yes"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[829] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[865] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[901] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[937] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[973] "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
      "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No" "No"
[ reached getOption("max.print") -- omitted 4000 entries ]
```

**iv. Compute the validation set error, which is the fraction of the observations in the validation set that are misclassified.**

```
> mean(glm_pred != Default[-train, ]$default)
[1] 0.0254
```

**(c) Repeat the process in (b) three times, using three different splits of the observations into a training set and a validation set. Comment on the results obtained.**

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit_glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> probs <- predict(fit_glm, newdata = Default[-train, ], type = "response")
> pred_glm <- rep("No", length(probs))
> pred_glm[probs > 0.5] <- "Yes"
> mean(pred_glm != Default[-train, ]$default)
[1] 0.0274
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit_glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> probs <- predict(fit_glm, newdata = Default[-train, ], type = "response")
> pred_glm <- rep("No", length(probs))
> pred_glm[probs > 0.5] <- "Yes"
> mean(pred_glm != Default[-train, ]$default)
```

```
[1] 0.0244
```

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit_glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> probs <- predict(fit_glm, newdata = Default[-train, ], type = "response")
> pred_glm <- rep("No", length(probs))
> pred_glm[probs > 0.5] <- "Yes"
> mean(pred_glm != Default[-train, ]$default)
```

```
[1] 0.0244
```

⇒ *The test error rate estimate changes with the observations included in training set and validation set.*

**(d) Consider another logistic regression model that predicts **default** using **income**, **balance** and **student** (qualitative). Estimate the test error for this model using the validation set approach. Does including the qualitative variable **student** lead to a reduction of test error rate?**

```
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit_glm <- glm(default ~ income + balance + student, data = Default, family = "binomial", subset = train)
> pred_glm <- rep("No", length(probs))
> probs <- predict(fit_glm, newdata = Default[-train, ], type = "response")
> pred_glm[probs > 0.5] <- "Yes"
> mean(pred_glm != Default[-train, ]$default)
```

```
[1] 0.0278
```

⇒ Yes, when we add the variable Student which is qualitative, the validation set test error rate decreases.

## Problem 2

**This question requires performing cross validation on a simulated data set.**

**(a) Generate a simulated data set as follows:**

```
set.seed(1)
x=rnorm(200)
y=x-2*x^2+rnorm(200)
```

**In this data set, what is  $n$  and what is  $p$ ? Write out the model used to generate the data in equation form (i.e., the true model of the data).**

```
> set.seed(4)
> x=rnorm(200)
> y=x-2*x^2+rnorm(200)
```

```
> p = 2
```

```
> n = 200
```

```
> simdata = data.frame(x,y)
```

```
> simdata
```

	x	y
1	0.2167548629	1.33751967
2	-0.5424925723	-2.67888922
3	0.8911446451	-0.99937894
4	0.5959805772	0.92480260
5	1.6356180011	-4.48271622
6	0.6892754419	1.26374675
7	-1.2812466301	-6.98651979
8	-0.2131445193	0.25232263
9	1.8965398719	-4.19165319
10	1.7768632137	-4.37118573
11	0.5666044982	-0.30093911
12	0.0157194540	-0.21318708
13	0.3830573385	-0.16359774
14	-0.0451371159	2.01906150
15	0.0343519074	1.61518166
16	0.1690267742	-0.93070407
17	1.1650268390	-1.55793563
18	-0.0442039973	-1.38219988
19	-0.1003684426	0.02601994
20	-0.2834445689	-1.23194964
21	1.5408149809	-3.49640423
22	0.1651690197	0.77744671
23	1.3076223603	-2.24876133
24	1.2882568779	-1.80686538
25	0.5928969406	1.00705004
26	-0.2829436843	-1.88155234
27	1.2558840256	-1.26423962
28	0.9098391512	-1.24497021
29	-0.9280281051	-1.52422800
30	1.2401808380	-1.93610087

31 0.1534641796 -1.00348035  
32 1.0519325790 -0.50326746  
33 -0.7542112128 -1.93517633  
34 -1.4821891197 -5.24756907  
35 0.8611318725 -1.55879354  
36 -0.4045198308 -1.08917582  
37 -0.2274054173 -0.37590478  
38 0.9340961709 -1.15902900  
39 -0.4658958798 -1.25898773  
40 -0.6375434986 -1.05485168  
41 1.3437086262 -3.25134128  
42 0.1815353846 0.09292463  
43 1.2925123364 -2.93321415  
44 -1.6880485759 -6.90202413  
45 -0.8209935776 -3.00058511  
46 -0.8621461441 -1.17558727  
47 0.0988436891 0.14031900  
48 -0.3756551442 -0.50113510  
49 0.7239041553 0.14473546  
50 -1.7973820186 -7.71338546  
51 -0.6637431416 -1.14925089  
52 -0.6237264887 -2.32315930  
53 -0.0796324318 -0.76888769  
54 0.4356247628 -1.88780841  
55 1.9709009697 -5.67801714  
56 -0.5967586725 -0.04292323  
57 -0.5525072116 -2.26102245  
58 0.6959666337 0.16513398  
59 -0.1556639646 0.41226654  
60 1.3488981952 -2.27228484  
61 -1.0685230705 -4.59119521  
62 1.0644507468 -1.70588097  
63 -1.3127217645 -3.44241695  
64 2.0636947023 -5.65954675  
65 0.1313830107 0.58189420

66	-0.2316884489	2.03198429
67	-0.3973555230	-0.99484339
68	0.8894320823	-0.49375130
69	0.5261690395	0.78269654
70	-0.1712732430	-0.35220088
71	0.1586768974	1.11603962
72	-0.4856650662	-0.43982404
73	-0.9589060750	-3.02787919
74	0.1805172921	-0.05228947
75	0.7217342828	-2.79610810
76	-0.3695404781	-1.24381503
77	0.2375383125	0.15619802
78	-0.6659221124	-0.32670204
79	-0.7968075098	-2.76582673
80	-0.0516969313	-1.56119087
81	1.2869283333	-2.25062174
82	-0.2141496627	-0.31480651
83	-0.5747454643	-2.71377182
84	-1.4707270443	-7.64455205
85	-1.0327384328	-4.01186658
86	-1.3065248552	-3.46479151
87	-0.8382524073	-3.59408856
88	-1.1306536810	-3.36359404
89	0.3687481753	0.21559276
90	-0.2018030203	0.36636424
91	-1.2776599028	-3.86087185
92	-0.7980124807	0.44585533
93	0.1590824229	0.41099298
94	0.6147976331	-0.80496528
95	0.6879479624	-0.19564098
96	-0.0470510111	-1.23914143
97	2.3303216783	-8.95097997
98	-0.5775659910	-1.99235139
99	0.9684791343	-1.03375179
100	-0.2775356275	-1.45025477


101	0.6848019360	-3.09268851
102	-0.1151135095	1.36985331
103	-0.3564751798	0.83173028
104	-0.1057716076	0.37722241
105	0.0448827901	-0.71643939
106	-1.7261732320	-7.43628320
107	1.5557870203	-3.60201208
108	0.7764126917	-0.32629930
109	-1.0985075088	-3.72355360
110	-1.7280197536	-7.42958121
111	0.4276382246	-0.24511093
112	0.7445646452	0.52568034
113	0.8652207970	-0.24394978
114	0.3053288101	-2.15551072
115	-0.1140227912	-0.68997687
116	0.4236522402	0.39506982
117	-0.7977096868	-2.66853581
118	-0.6041972494	-2.59635014
119	1.7150105938	-3.40421367
120	-0.7159482778	-0.90014751
121	-0.1332356122	-0.01051404
122	-0.9997650626	-1.95232731
123	1.8737601171	-5.80104647
124	-0.3373884320	-1.52577030
125	0.9732702887	-1.12238823
126	0.9878279309	-0.27019526
127	-0.9412566085	-3.71121148
128	0.3491855939	-0.89431238
129	-0.5944186817	-1.66440068
130	-2.3822428313	-12.67473814
131	1.0780189737	-0.16727267
132	0.6682451050	1.01313897
133	-0.9646256667	-1.97151861
134	-1.9752373319	-10.59190125
135	-0.5847739007	-1.63653000



136	0.9692770362	0.20796480
137	0.5522923259	0.05596528
138	-0.0821555007	-0.54019582
139	-1.6767137584	-8.74517167
140	1.2126074270	-2.11267032
141	1.0004998710	-0.40837191
142	0.7193289908	-0.07048793
143	-0.8443641520	-4.01606610
144	0.6219853903	-0.08932241
145	-0.7226137804	-2.85298784
146	-0.4494786251	-0.96106433
147	-1.1955060501	-3.60278091
148	0.3904723630	1.17990462
149	-0.5163766426	-0.14370718
150	0.9098689779	-1.74336520
151	0.8769846530	-1.37705875
152	-0.8161958099	-1.72594614
153	1.5392932699	-2.76662090
154	1.3745257156	-2.31958453
155	-0.4832487112	-0.39821124
156	0.5503499503	-0.59318566
157	-0.8573656630	-2.20177086
158	-0.7069613662	-0.70266279
159	-2.0970775334	-9.82070865
160	1.0994367548	-1.76843817
161	0.3420340890	1.06650105
162	0.4908294804	0.91332704
163	-0.9319990260	-3.90057416
164	-1.4278919839	-5.79609941
165	0.9757650946	-2.58057887
166	-1.5463411878	-6.29127356
167	0.0177034792	-2.02560374
168	-0.7747174012	-1.86382201
169	-0.2293422872	-0.02333439
170	-0.2743821044	-1.06385761

```
171 1.7960637815 -3.72286011
172 -0.4781128994 -1.20797494
173 -0.5947628530 -0.06040863
174 -2.2579382170 -13.51519417
175 1.6826072118 -2.09161048
176 0.0722906844 -1.10625306
177 -0.4400240932 -0.40826205
178 0.6265733926 -1.81109669
179 -0.7997960594 -1.90223804
180 -1.1279860222 -3.81661442
181 -1.0250160534 -0.76362512
182 0.0710717295 1.27793900
183 0.3817111616 0.49359057
184 -1.6225883175 -6.25054679
185 1.9005426699 -6.21267209
186 -0.7161791664 -3.05721269
187 0.3804596689 0.26909412
188 0.4408428474 -1.37164919
189 0.2573258583 -1.37879719
190 -0.1794485371 -1.31925064
191 -0.6901276793 -1.50088908
192 -0.0004228025 0.26177169
193 0.5655808964 1.11973408
194 -1.2087470098 -5.81777880
195 -0.3461711560 -0.34579272
196 -0.6501970444 -2.62023640
197 -0.8895916708 -3.32253969
198 1.4770298873 -3.22494027
199 -1.1954751385 -4.11350783
200 1.7504948348 -4.53185060
> dim(simdata)
[1] 200 2
```

(b) Create a scatter plot of  $Y$  vs  $X$ . Comment on what you find.



$\Rightarrow$  From the plot obtained above, the data looks like a parabolic distribution. (atleast approximately)

(c) Consider the following four models for the data set:

i.  $Y = \beta_0 + \beta_1 X + \epsilon$

ii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \epsilon$

iii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$

iv.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 X^4 + \epsilon$

Compute the LOOCV errors that result from fitting these models.

```
> library(boot)
> glm.fit = glm(y~x,data=simdata)
> cv.err = cv.glm(simdata,glm.fit)
> cv.error = rep(0,4)
> for (i in 1:4){
  glm.fit = glm(y~poly(x,i),data=simdata)
  cv.error[i] = cv.glm(simdata,glm.fit)$delta[1]
}
> cv.error
[1] 6.0292645 0.9487129 0.9597797 0.9742622
```

$\Rightarrow$  LOOCV can be computed directly for generalized linear models using `glm()` and `cv.glm()`.

⇒ Since, linear reg belongs to generalized linear models, we can use the glm() rather than lm() to fit the model.

**(d) Repeat (c) using another random seed, and report your results. Are your results the same as what you got in (c)? Why?**

```
> set.seed(2)
```

```
> x1=rnorm(200)
```

```
> y1=x1-2*x1^2+rnorm(200)
```

```
> simdata1 = data.frame(x1,y1)
```

	x1	y1
1	-0.896914547	-2.207842334
2	0.184849185	-0.903041455
3	1.587845331	-0.583762837
4	-1.130375674	-3.467163993
5	-0.080251757	-1.059686775
6	0.132420284	0.481188203
7	0.707954729	-0.416341139
8	-0.239698024	-0.705155882
9	1.984473937	-5.292336053
10	-0.138787012	0.053963396
11	0.417650751	1.095032648
12	0.981752777	-1.470373105
13	-0.392695356	1.095091602
14	-1.039668977	-4.635688546
15	1.782228960	-4.431765003
16	-2.311069085	-12.541685985
17	0.878604581	0.549486686
18	0.035806718	-1.291238649
19	1.012828692	-2.178443003
20	0.432265155	1.734207397
21	2.090819205	-6.254084484
22	-1.199925820	-3.364050630
23	1.589638200	-4.257456151
24	1.954651642	-3.762822796
25	0.004937777	0.075329000
26	-2.451706388	-14.004154776

27	0.477237303	-0.068125677
28	-0.596558169	-0.414169834
29	0.792203270	-0.836678651
30	0.289636710	2.380109766
31	0.738938604	1.219979864
32	0.318960401	-1.838832805
33	1.076164354	-1.859417513
34	-0.284157720	-0.475596669
35	-0.776675274	-3.704782658
36	-0.595660499	-2.358740157
37	-1.725979779	-8.254654659
38	-0.902584480	-2.049187571
39	-0.559061915	-1.088238472
40	-0.246512567	-0.955889250
41	-0.383586228	-0.759643054
42	-1.959103175	-6.746855423
43	-0.841705060	-1.085814524
44	1.903547467	-4.793795233
45	0.622493930	0.626601054
46	1.990920436	-6.501739177
47	-0.305483725	-0.389170051
48	-0.090844235	0.174171928
49	-0.184161452	-0.726423506
50	-1.198767765	-4.530860283
51	-0.838287148	-1.742194659
52	2.066301356	-6.878030606
53	-0.562247053	-1.371876142
54	1.275715512	-2.298729192
55	-1.047572627	-2.398939052
56	-1.965878241	-9.211576113
57	-0.322971094	-0.835261896
58	0.935862527	-1.100910908
59	1.139229803	-0.476022607
60	1.671618767	-4.088319308
61	-1.788242207	-6.840849445


62	2.031242519	-6.319008256
63	-0.703144333	-2.013963719
64	0.158164763	-1.806601923
65	0.506234797	-0.653353936
66	-0.819995106	-3.564215899
67	-1.998846995	-8.058451344
68	-0.479292591	-0.477163075
69	0.084179904	0.597446138
70	-0.895486611	-2.535970423
71	-0.921275666	-5.338941397
72	0.330449503	0.064930239
73	-0.141660809	-0.177059907
74	0.434847762	-1.168641133
75	-0.053722626	0.191430348
76	-0.907110376	-0.947608906
77	1.303512232	-2.130295925
78	0.771789776	0.609888951
79	1.052525595	-1.563137581
80	-1.410038341	-5.319211370
81	0.995984590	1.047821839
82	-1.695764903	-6.797483312
83	-0.533372143	0.177007703
84	-1.372269451	-5.766715681
85	-2.207919779	-10.607407255
86	1.822122519	-3.145548452
87	-0.653393411	-0.334509752
88	-0.284681219	-0.440589076
89	-0.386949604	0.619102625
90	0.386694975	-0.005357864
91	1.600390852	-4.650860720
92	1.681154956	-3.384062630
93	-1.183606388	-3.896290671
94	-1.358457254	-5.282233898
95	-1.512670795	-4.601240782
96	-1.253104899	-4.747741938

97 1.959357077 -5.300154809  
98 0.007645872 -2.089906587  
99 -0.842615198 -3.633096121  
100 -0.601160105 -2.008409906  
101 1.074459406 -1.552586468  
102 0.260597835 -0.190714883  
103 -0.314271980 0.372516588  
104 -0.749630095 -3.758941961  
105 -0.862198330 -1.616790921  
106 2.048040303 -5.550353199  
107 0.939920078 -0.632126263  
108 2.008687116 -5.139782751  
109 -0.421373572 0.464971260  
110 -0.350834423 -1.278174810  
111 -1.027380598 -2.623071921  
112 -0.250519127 -0.844249877  
113 0.471859466 -0.475500567  
114 1.358939821 -4.308486928  
115 0.564168603 -0.641849716  
116 0.455980090 0.515474958  
117 1.230953663 -0.582322489  
118 1.147136848 -1.367019479  
119 0.106598041 0.172238471  
120 -0.783316657 -2.211103708  
121 1.241199827 -2.338537751  
122 0.138858419 1.189962637  
123 1.710631588 -3.220258494  
124 -0.430640975 -0.436990715  
125 -1.044229581 -2.929912262  
126 0.537579525 0.489544452  
127 -0.669585987 -0.597486973  
128 0.638805611 1.731266130  
129 -1.723989834 -6.165478286  
130 -1.742430080 -8.126383194  
131 0.689804173 0.686163264

132	0.330963177	0.024217166
133	0.871067709	-0.840176510
134	-2.016245582	-9.417497348
135	1.212579104	-0.699712203
136	1.200494699	-1.989070683
137	1.032068326	-0.332520138
138	0.786410256	-0.471922659
139	2.110073514	-5.302219068
140	-1.453809847	-5.233257915
141	-0.583103848	-1.088147214
142	0.409723983	0.461217718
143	-0.806981635	-3.398314798
144	0.085550441	0.302419518
145	0.746243169	0.352632287
146	-0.653673061	-1.759681926
147	0.657105983	-0.525466447
148	0.549909235	0.928552326
149	-0.806729358	-1.280112918
150	-0.997379717	-2.604118451
151	0.975890638	0.526825322
152	-0.169423181	0.439628900
153	0.722191779	-0.690937408
154	-0.844418607	-0.722252935
155	1.277293685	0.343599389
156	-1.343110549	-3.551577872
157	0.765340669	-1.347397830
158	0.464202570	0.866535366
159	0.267993278	0.859239076
160	0.667522687	-0.100432217
161	0.398467284	1.636278509
162	-0.638071031	-0.029082227
163	-0.267712904	-0.769810465
164	0.359879565	-0.396890260
165	-1.312866094	-6.645088891
166	-0.883969610	-3.586400435



```
167 2.077094795 -7.008279795
168 -2.099225632 -9.739388568
169 -1.238505965 -4.626487321
170 0.990433090 -0.929809062
171 1.088661863 -1.458690619
172 0.839852254 1.713886253
173 0.056858639 -0.753217649
174 0.323878051 1.344717967
175 -0.904668668 -2.575964043
176 -0.652183848 -2.165312257
177 -0.262454638 -0.649819693
178 -0.934662841 -0.148738020
179 0.821161213 -0.369330966
180 -1.624259173 -6.148655723
181 -1.030403669 -3.581591123
182 -1.261929312 -5.185608814
183 0.392184626 0.228214426
184 -1.131438262 -4.406765221
185 0.544144484 1.268443711
186 1.176608935 -3.166475546
187 0.025228569 0.070116689
188 0.515133170 -0.044135014
189 -0.654109760 -0.096700115
190 0.503641991 0.825524488
191 -1.272119222 -4.325341777
192 -0.076771154 1.250998232
193 -1.345319376 -3.762758692
194 -0.266317560 0.463667903
195 1.087562995 -1.394450370
196 0.700567795 0.061805803
197 -0.442759515 -1.820302827
198 -0.788519966 -0.904335964
199 -0.856775710 -2.095075092
200 -0.746419015 -0.942131153
> plot(simdata1)
```



```
> glm.fit1 = glm(y1~x1,data=simdata1)
> cv.err1 = cv.glm(simdata1,glm.fit1)
>
> cv.error1 = rep(0,4)
> for (i in 1:4){
  glm.fit1 = glm(y1~poly(x1,i),data=simdata1)
  cv.error1[i] = cv.glm(simdata1,glm.fit1)$delta[1]
}
> cv.error1
[1] 7.1017678 0.9627777 0.9633221 0.9761597
```

⇒ Yes, the results obtained are close to the previous because LOOCV works as whole dataset minus one observation. Again, the test error is smallest for second degree quadratic model.

**(e) Which of the models in (c) had the smallest LOOCV error? Is this what you expected? Explain your answer.**

⇒ Quadratic model ( $2^{ND}$  degree) has the smallest LOOCV error. This is to be expected because the data in original model is in quadratic form.

**(f) Now we use 5-fold CV for the model selection. Compute the CV errors that result from fitting the four models. Which model has the smallest CV error? Are the results consistent with LOOCV?**

```
> library(boot)
> set.seed(5)
> cv.error.5 = rep(0,5)
> for (i in 1:5) {
  glm.fit = glm(y~poly(x,i),data=simdata)
```

```

cv.error.5[i] = cv.glm(simdata,glm.fit,K=5)$delta[1]
}
> cv.error.5
[1] 6.1170694 0.9544515 0.9697376 0.9607885 0.9789105

```

⇒ CV error is again smallest for model with degree 2. These results are consistent with LOOCV.

**(g) Repeat (f) using 10-fold CV. Are the results the same as 5-fold CV?**

```

> library(boot)
> set.seed(2)
> cv.error.10 = rep(0,10)
> for (i in 1:10) {
  glm.fit = glm(mpg~poly(horsepower,i),data=Auto)
  cv.error.10[i] = cv.glm(Auto,glm.fit,K=10)$delta[1]
}
> cv.error.10
[1] 24.22275 19.14140 19.31435 19.33767 19.13637 18.95081 18.82974 19.50147 19.16699
19.19373

```

⇒ For 10-fold CV, model with degree 7 has the least CV error, unlike 5-fold CV where 2<sup>nd</sup> degree polynomial had least CV error.